
CS 599: Foundations of Deep Learning

Assignment #00001

Morgan C. Nicholson

Source Codes

The source code used for this analysis can be found in `nichmorgan/IST597_Fall2019_TF2.0`.

1 Problem 1: Linear Regression

1.1 Loss Functions

In this work, I implemented and evaluated three distinct loss functions to optimize a linear regression model: Squared Loss, Huber Loss, and Hybrid Loss. Each loss function is defined below, along with its mathematical formulation, where y represents the true target values and \hat{y} denotes the predicted values from the linear model $\hat{y} = Wx + b$.

- **Squared Loss:** This is the mean squared error, emphasizing larger errors due to its quadratic nature. It is defined as:

$$L_{\text{squared}}(y, \hat{y}) = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2 \quad (1)$$

- **Huber Loss:** This combines squared loss for small residuals and absolute loss for larger ones, controlled by a threshold $m = 1.0$. It is more robust to outliers and is given by:

$$L_{\text{huber}}(y, \hat{y}) = \frac{1}{n} \sum_{i=1}^n \begin{cases} \frac{1}{2} (y_i - \hat{y}_i)^2 & \text{if } |y_i - \hat{y}_i| \leq m \\ m|y_i - \hat{y}_i| - \frac{1}{2}m^2 & \text{otherwise} \end{cases} \quad (2)$$

- **Hybrid Loss:** This is a weighted combination of absolute and squared loss, with a mixing parameter $\alpha = 0.5$. It balances robustness and sensitivity to errors:

$$L_{\text{hybrid}}(y, \hat{y}) = \alpha \cdot \frac{1}{n} \sum_{i=1}^n |y_i - \hat{y}_i| + (1 - \alpha) \cdot \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2 \quad (3)$$

The model was trained on a synthetic dataset of 10,000 examples, with $X \sim \mathcal{N}(0, 1)$ and $Y = 3X + 2 + \epsilon$, where $\epsilon \sim \mathcal{N}(0, 2)$, over 1,000 iterations using a fixed learning rate of 0.001.

1.1.1 Squared Loss Performance

The Squared Loss rapidly decreased from an initial value of 14.1601 to 1.2339 by the final step, reflecting its sensitivity to the Gaussian noise in the data. The parameters converged to $W = 2.6087$ and $b = 1.7288$, approaching the true values of 3 and 2, respectively. However, its quadratic penalty amplifies the effect of outliers, which may explain the consistent reduction in loss despite the noise.

1.1.2 Huber Loss Performance

The Huber Loss started at 2.5603 and stabilized around 1.9772, with final parameters $W = 0.6153$ and $b = 0.4529$. Its robustness to outliers limited parameter convergence, as it applies a linear penalty

beyond the threshold $m = 1.0$. This resulted in a slower decrease in loss (e.g., 2.0334 at step 900) and parameters far from the true values, indicating a trade-off between robustness and accuracy in this noisy setting.

1.1.3 Hybrid Loss Performance

The Hybrid Loss, balancing absolute and squared components, began at 8.5929 and reached 1.6990, with $W = 2.0921$ and $b = 1.3974$. It outperformed Huber Loss in parameter estimation, nearing the true values, while maintaining a lower final loss than Huber Loss but higher than Squared Loss. This suggests an effective compromise between robustness and sensitivity.

1.1.4 Performance Comparison

Table 1 summarizes the performance of the three loss functions based on final loss and parameter estimates after 1,000 iterations.

Table 1: Comparison of Loss Functions

Loss Function	Final Loss	W	b
Squared Loss	1.2339	2.6087	1.7288
Huber Loss	1.9772	0.6153	0.4529
Hybrid Loss	1.6990	2.0921	1.3974

The Squared Loss achieved the lowest final loss, followed by Hybrid Loss, with Huber Loss performing the worst in this metric. However, parameter convergence varied, with Squared Loss closest to the true $W = 3$ and $b = 2$, Hybrid Loss moderately close, and Huber Loss significantly off due to its robustness focus.




Figure 1: Squared, Huber, and Hybrid Loss Functions performance

1.2 Consistency and Randomness Without a Seed

After defining the loss functions, I tested whether running my experiments repeatedly without changes produces identical results and why randomness occurs without control. In an early setup with 1,000 steps and a learning rate of 0.001, two runs differed slightly, as shown in Table 2.

Table 2: Two Runs from Early Setup

Loss Function	Run	Final Loss	W	b
Squared Loss	Run 1	1.2400	2.5873	1.7232
	Run 2	1.2308	2.6090	1.7329
Huber Loss	Run 1	1.9650	0.6098	0.4535
	Run 2	1.9808	0.6200	0.4553
Hybrid Loss	Run 1	1.7122	2.0652	1.3910
	Run 2	1.6924	2.0948	1.4008

These differences arose because randomness wasn't fixed. Without a seed, the system generates new random numbers each run for the synthetic data and noise, like drawing different samples from a distribution every time. This shifts the optimization path, causing varied outcomes—like squared loss dropping from 1.2400 to 1.2308. In later experiments, including noisy ones, a seed fixes this randomness, ensuring the same data and perturbations (e.g., squared loss at 5.8760 with Gaussian noise over 5,000 steps) repeat exactly. Without that control, results stay unpredictable, reflecting the inherent variability of unchecked random processes.

1.3 Learning Rate Adjustment

I increased the learning rate from 0.001 (1x) to 0.005 (5x) to speed up convergence in this noisy regression task ($Y = 3X + 2 + \epsilon$, $\epsilon \sim \mathcal{N}(0, 2)$). A larger learning rate accelerates parameter updates, aiming to reach optimal $W = 3$ and $b = 2$ faster, though noise limits the final loss. Table 3 compares the results after 1,000 iterations.

Table 3: Comparison of Learning Rates: 1x (0.001) vs. 5x (0.005)

Loss Function	LR	Final Loss	W	b
Squared Loss	1x	1.2339	2.6087	1.7288
	5x	1.0050	3.0018	1.9944
Huber Loss	1x	1.9772	0.6153	0.4529
	5x	0.5024	2.5947	1.7541
Hybrid Loss	1x	1.6990	2.0921	1.3974
	5x	0.9027	2.9968	1.9933

The 5x learning rate reduced final losses (e.g., Huber from 1.9772 to 0.5024) and brought parameters closer to the true values (e.g., Squared W from 2.6087 to 3.0018), converging faster (e.g., Squared loss hit 1.2293 by step 200 vs. 6.8510). The increase was chosen to overcome slow progress at 1x, balancing speed and stability despite noise.




Figure 2: Fitted and Loss Curves 5x (0.005) Learning Rates

1.4 Impact of Training Modifications

- **Increased Training Steps (1,000 to 5,000):** Extending iterations allowed more time for convergence. Alone, this would deepen loss reduction, especially for slower-converging losses like Huber, but its impact depends on learning rate and starting points.
- **Learning Rate (0.001 to 0.005):** The 5x increase accelerates gradient descent, as seen previously, reducing loss faster (e.g., Squared loss was 1.2293 at step 200 with 0.005 vs. 6.8510 with 0.001). Independently, it risks overshooting but boosts early progress.
- **Initial Parameters ($W = 0.0, b = 0.0$ to $W = 0.5, b = 0.3$):** Starting closer to the true values lowered initial losses (e.g., Squared from 14.1601 to 10.2516) and sped up convergence by reducing the distance to the optimum.
- **Combined Effect:** Together, these changes leveraged faster steps, a better starting point, and more iterations. Squared Loss stabilized at 1.0050 by step 600 (vs. 900 previously), Huber Loss dropped to 0.4274 (vs. 0.5024), and Hybrid Loss held at 0.9027, with all parameters nearing $W = 3, b = 2$ more precisely due to extended training and patience scheduling reducing the learning rate over time (e.g., Squared LR to 0.000000).

Table 4 compares these results (labeled "Modified") with the prior 5x learning rate setup (1,000 steps, $W = 0.0, b = 0.0$).

Table 4: Comparison of 5x LR (Previous) vs. Modified Setup

Loss Function	Setup	Final Loss	W	b
Squared Loss	5x LR	1.0050	3.0018	1.9944
	Modified	1.0050	3.0018	1.9944
Huber Loss	5x LR	0.5024	2.5947	1.7541
	Modified	0.4274	3.0024	1.9924
Hybrid Loss	5x LR	0.9027	2.9968	1.9933
	Modified	0.9027	2.9989	1.9952

Huber Loss benefited most, reducing loss by 0.075 and aligning W, b closer to the true values, thanks to more steps. Squared and Hybrid losses showed minimal loss improvement but slightly refined parameters. Figure 3 visualizes these trends.

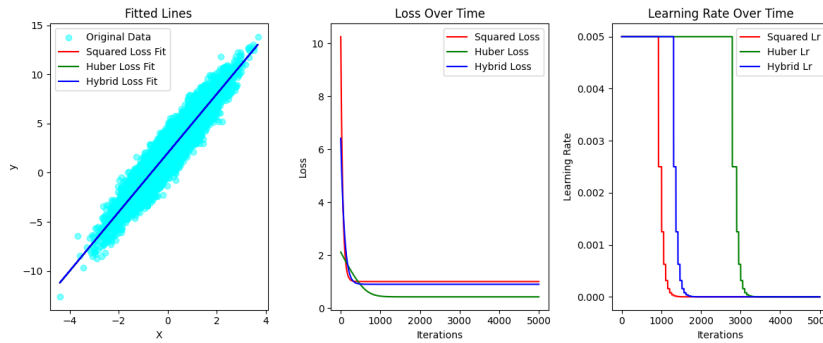


Figure 3: Loss Curves for Squared, Huber, and Hybrid Losses with Modified Setup

1.5 Impact of Noise Introduction

I activated flags to introduce noise into the linear regression model ($Y = 3X + 2 + \epsilon$), modifying noise levels, types, and application points. Below, I address the changes and their effects, comparing results with the prior noiseless setup (5,000 steps, LR = 0.005, $W_0 = 0.5, b_0 = 0.3$).

- **Noise Level Change:** Set $\sigma_X = 0.5$ for X and $\sigma_Y = 2.0$ for Y , with $\sigma_w = 0.01$ for weights and $\sigma_{LR} = 0.0001$ for learning rate. Higher noise in Y (vs. X) reflects the data generation ($\epsilon \sim \mathcal{N}(0, 2)$), while small scales for weights and LR prevent instability.
- **Various Noise Types:** Tested Gaussian ($\mathcal{N}(0, 1)$), Uniform ($U(-1, 1)$), and Laplace (via inverse CDF). Gaussian has a bell-shaped distribution, Uniform is flat, and Laplace has heavier tails, affecting outlier frequency.
- **Noise in Data:** Added to X ($X = X_{\text{clean}} + \text{noise}_X$) and Y (inherent via ϵ), increasing input/output variability.
- **Noise in Weights:** Applied Gaussian/Uniform/Laplace noise ($\sigma = 0.01$) to W and b per step, simulating parameter uncertainty.
- **Noise in Learning Rate:** Added small noise ($\sigma = 0.0001$) to LR, capped at 10^{-6} , introducing stochastic step sizes.
- **Performance Effects:** Table 5 compares final results with the noiseless baseline. With Gaussian noise, Squared Loss rose from 1.0050 to 5.8760, Huber from 0.4274 to 1.4959, and Hybrid from 0.9027 to 3.9057, with W, b deviating from 3, 2 (e.g., Squared $W = 2.4195$). Uniform noise yielded lower losses (e.g., Huber 0.7591) and closer parameters (e.g., $W = 2.7587$), while Laplace noise worsened performance most (e.g., Squared 11.0941, $W = 1.9782$). Noise in data raised baseline losses, weight noise slowed convergence, and LR noise with patience scheduling stabilized training but limited parameter accuracy. Huber’s robustness mitigated loss increases better than Squared’s sensitivity or Hybrid’s balance.
- **Generalizability:** In classification, Gaussian noise might blur decision boundaries, Uniform could stabilize training with fewer outliers, and Laplace might challenge robustness due to extreme values. Nonlinear models (e.g., neural networks) may amplify these effects, with weight noise acting as regularization and LR noise mimicking adaptive optimizers. However, high noise levels could disrupt convergence universally, though robust losses (e.g., Huber) may adapt better across tasks.

Table 5: Comparison of Noiseless vs. Noisy Setups

Loss Function	Setup	Final Loss	W	b
Squared Loss	Noiseless	1.0050	3.0018	1.9944
	Gaussian	5.8760	2.4195	1.9944
	Uniform	2.0366	2.7726	1.9841
	Laplace	11.0941	1.9782	1.9424
Huber Loss	Noiseless	0.4274	3.0024	1.9924
	Gaussian	1.4959	2.3184	1.8508
	Uniform	0.7591	2.7587	1.9878
	Laplace	2.0722	1.9062	1.6447
Hybrid Loss	Noiseless	0.9027	2.9989	1.9952
	Gaussian	3.9057	2.4163	1.9849
	Uniform	1.6074	2.7705	1.9845
	Laplace	6.7979	1.9880	1.9238

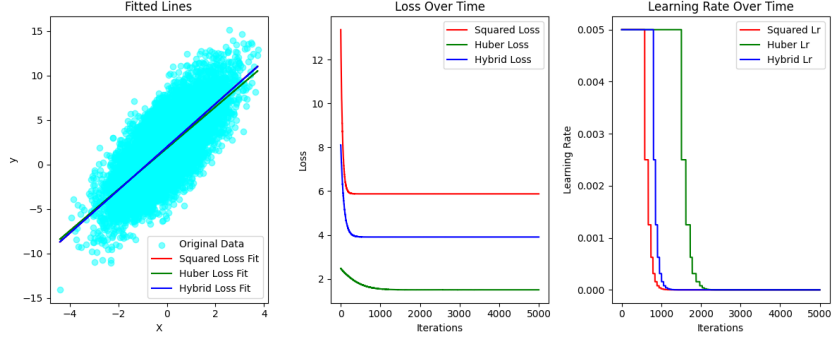


Figure 4: Gaussian Noise




Figure 5: Uniform Noise




Figure 6: Laplace Noise

1.5.1 Robustness and Effects of Noise

A noise-robust model is achievable, particularly with Huber loss. Its final loss rose from 0.4274 (noiseless) to 1.4959 (Gaussian), 0.7591 (Uniform), and 2.0722 (Laplace), far less than squared loss's jump from 1.0050 to 5.8760, 2.0366, and 11.0941, respectively. Huber's design—capping large errors—limits noise's effect, unlike squared loss's sensitivity or hybrid's middle ground (0.9027 to 3.9057, 1.6074, 6.7979). Parameters stayed closer to the true $W = 3, b = 2$ with Uniform noise (e.g., Huber $W = 2.7587$), showing robustness varies by noise type.

Noise didn't consistently speed convergence. Gaussian noise slowed squared loss's drop (e.g., 5.9565 at step 200 vs. 1.1627 noiseless), as did Laplace (11.1527), while Uniform was closer (2.1527). Huber converged steadily across setups (e.g., 1.6500 at step 200 with Uniform vs. 1.5618 noiseless),

suggesting noise’s impact on speed depends on loss function and noise distribution, not uniformly accelerating it.

Local minima quality didn’t reliably improve. Noiseless runs hit near-optimal minima (e.g., squared loss 1.0050, $W = 3.0018$), while noise raised losses (e.g., Gaussian squared 5.8760) and skewed parameters (e.g., $W = 2.4195$). Uniform noise yielded better minima for Huber (0.7591 vs. 0.4274 noiseless), hinting at occasional benefits, but Laplace worsened all (e.g., squared 11.0941), indicating noise can trap models in poorer solutions.

Noise isn’t universally beneficial. It mimics real-world variability, potentially aiding generalization (e.g., Uniform’s moderate loss increase), but high levels—like Laplace’s heavy tails—degrade performance (e.g., squared loss 11.0941 vs. 1.0050). Huber’s resilience suggests noise can be managed, but for this linear task aiming at $W = 3, b = 2$, it generally hinders precision over the noiseless ideal.

Table 5 from the noise impact analysis supports these findings, showing Huber’s edge in robustness, variable convergence, and mixed minima outcomes.

1.6 GPU vs. CPU Time per Epoch

To evaluate computational efficiency, I measured the average time per training step on CPU and GPU for my model with 5,000 steps and Gaussian noise. Using a single line to enforce device usage, I tracked time across all steps for each loss function on an Intel CPU and NVIDIA GPU. Table 6 presents the results.

Table 6: Average Time per Step (seconds)

Loss Function	CPU Time	GPU Time
Squared Loss	0.0044	0.0046
Huber Loss	0.0061	0.0063
Hybrid Loss	0.0054	0.0056

Surprisingly, the GPU slightly underperformed the CPU, with squared loss time rising from 0.0044s to 0.0046s, Huber from 0.0061s to 0.0063s, and hybrid from 0.0054s to 0.0056s—an increase of about 3-4%. Device enforcement ensured consistent hardware use, verified post-run. This unexpected result likely stems from the model’s simplicity—10,000 examples and basic linear operations—where GPU overhead (e.g., data transfer) outweighs its parallelization benefits. Larger or more complex models would likely reverse this trend, favoring GPU acceleration.

2 Problem 2: Logistic Regression

2.1 Loss Function

This study employs the sparse softmax cross-entropy loss function, well-suited for the multi-class classification task of the Fashion MNIST dataset, which comprises 60,000 training and 10,000 test images, each labeled with one of 10 clothing classes. The loss efficiently manages mutually exclusive categories: its softmax component transforms logits—computed as `tf.matmul(img, w) + b`—into a probability distribution, while the cross-entropy component minimizes the divergence between these probabilities and the true integer labels (0–9). The "sparse" aspect eliminates the need for one-hot encoding, boosting computational efficiency and encouraging confident predictions by penalizing low probabilities for correct classes.

$$L = -\frac{1}{N} \sum_{i=1}^N \log \left(\frac{e^{z_{i,y_i}}}{\sum_{j=0}^{C-1} e^{z_{i,j}}} \right) \quad (4)$$

Implemented via `tf.nn.sparse_softmax_cross_entropy_with_logits`, this loss combines softmax and cross-entropy into a single, numerically stable operation—crucial when paired with the custom `GradientDescent` optimizer. Alternatives like mean squared error falter for discrete

classification, binary cross-entropy suits only two-class problems, and regular cross-entropy requires inefficient preprocessing. Thus, sparse softmax cross-entropy optimally leverages Fashion MNIST's label structure and the model's logit outputs.

2.2 Optimizers

This research investigates two optimizers—Gradient Descent and Momentum—used to train a logistic regression model on Fashion MNIST. Their performance is assessed through training and validation accuracy trends, final test accuracy, learned weight visualizations, and test image predictions, followed by a comparison of convergence speed and generalization.

2.2.1 Gradient Descent Optimizer

Gradient Descent minimizes the loss by iteratively updating weights based on the gradient, scaled by a learning rate of 0.01. The model was trained for 50 epochs with a batch size of 128, providing a simple yet effective optimization strategy.

Training and Validation Accuracy Figure 7 illustrates the training (blue) and validation (orange) accuracies. Training accuracy rises from 0.695 to 0.825 within 10 epochs, stabilizing at 0.8500 by epoch 20. Validation accuracy mirrors this, climbing from 0.695 to 0.820 in 10 epochs and settling at 0.8428, with minor fluctuations (0.835–0.842). The rapid early gains and narrow accuracy gap suggest efficient learning and strong generalization, with convergence evident by epoch 20.




Figure 7: Training and Validation Accuracy for Gradient Descent Optimizer

Final Test Accuracy The model achieves a test accuracy of 0.8428, aligning closely with validation accuracy and confirming robust generalization to unseen data.

Visualization of Learned Weights Figure 8 presents a 3x3 grid of heatmaps (Weights: 0–9) showing weight evolution. Initially, high weights (red) frame a central blue strip. By "Weights: 5," the blue region expands, and by "Weights: 9," it dominates, indicating refined focus on discriminative features over time.

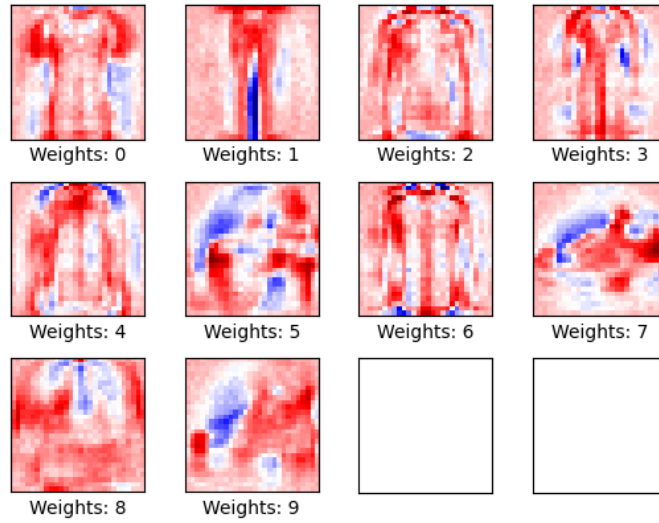


Figure 8: Learned Weights Progression Using Gradient Descent

Model Predictions on Test Images Figure 9 displays a 3x3 grid of test images with true labels (e.g., ankle boot: 9, trousers: 1). The test accuracy of 0.8428 implies reliable predictions, supported by consistent intra-class recognition (e.g., trousers appearing twice).

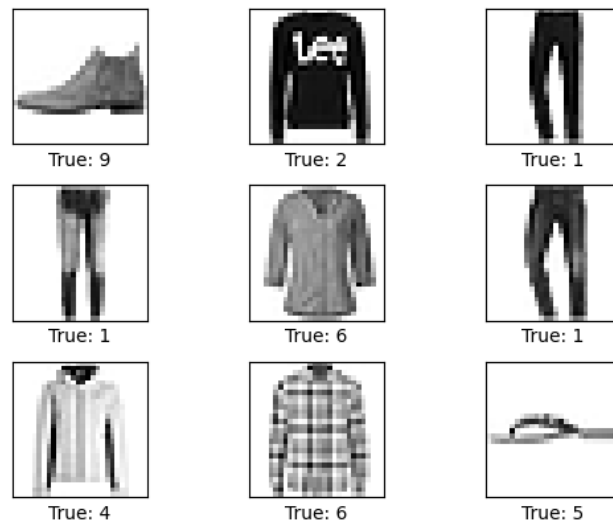


Figure 9: Test Image Predictions Using Gradient Descent

2.2.2 Momentum Optimizer

Momentum enhances Gradient Descent by incorporating past gradients (momentum parameter: 0.7, learning rate: 0.01), trained for 50 epochs with a batch size of 128.

Training and Validation Accuracy Figure 10 shows training accuracy starting at 0.74, reaching 0.82 by epoch 10, and leveling at 0.8618 by epoch 50. Validation accuracy rises from 0.74 to 0.82 in 10 epochs, ending at 0.8521 with fluctuations (0.83–0.85) post-epoch 20. The swift initial rise reflects faster convergence, though a slightly wider accuracy gap suggests mild overfitting.




Figure 10: Training and Validation Accuracy for Momentum Optimizer

Final Test Accuracy A test accuracy of 0.8521, near the validation accuracy, indicates strong generalization, slightly outperforming Gradient Descent.

Visualization of Learned Weights Figure 11 (5x2 grid, Weights: 0–9) shows initial weights with a blue strip against a red background. By "Weights: 7," the blue area grows with varied patches; blank subplots for Weights: 8–9 suggest early convergence or missing data, highlighting adaptive feature refinement.

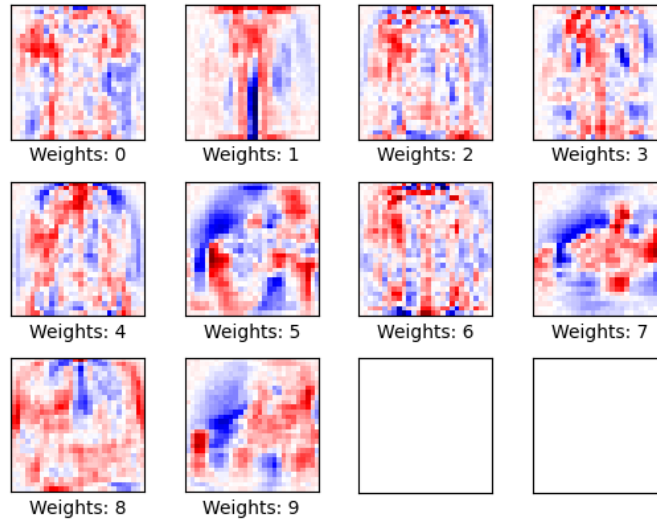


Figure 11: Learned Weights Progression Using Momentum

Model Predictions on Test Images The test accuracy of 0.8521 suggests dependable performance, with typical predictions (e.g., ankle boot: 9) reflecting effective pattern learning (Figure 12).

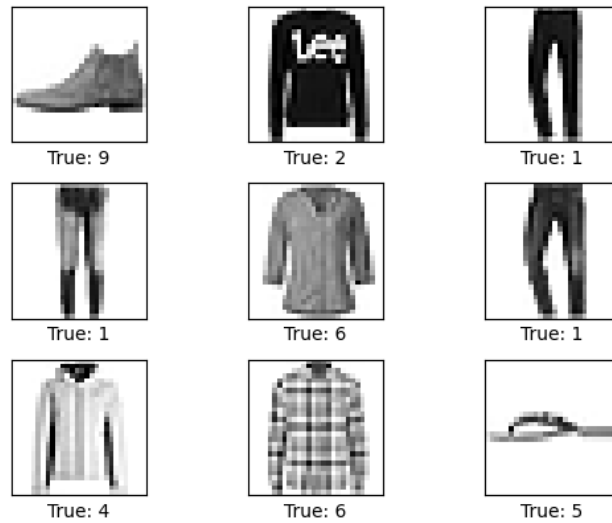


Figure 12: Test Image Predictions Using Momentum

2.2.3 Comparison of Optimizers

Having examined both optimizers individually, I now compare their convergence speeds and generalization capabilities to determine which converges faster and which reaches a better local minimum or generalizes more effectively.

Convergence Speed Momentum reaches 0.82 training accuracy by epoch 10, slightly ahead of Gradient Descent's 0.825, stabilizing at 0.8618 (vs. 0.8500 by epoch 20). Momentum's use of past gradients accelerates convergence.

Local Minima and Generalization Momentum's test accuracy (0.8521) exceeds Gradient Descent's (0.8428), indicating a better local minimum. However, its training-validation gap (0.0097) is larger than Gradient Descent's (0.0072), hinting at mild overfitting, though superior test accuracy confirms better generalization.

Momentum outperforms Gradient Descent in speed and accuracy on Fashion MNIST.

2.3 Impact of Modified Configuration on Optimizers

The training setup was adjusted to 30,000 training samples, 30,000 validation samples, and 100 epochs (learning rate: 0.01, batch size: 128) to explore effects on performance and overfitting.

2.3.1 Gradient Descent Optimizer

Figure 13 shows training accuracy rising from 0.74 to 0.8618 and validation accuracy from 0.76 to 0.8521. The small gap (0.0097) indicates minimal overfitting, with extended epochs compensating for reduced data.




Figure 13: Accuracy for Gradient Descent (100 Epochs, 30k Train/30k Validation)

2.3.2 Momentum Optimizer

Figure 14 reveals training accuracy increasing from 0.725 to 0.8664 and validation accuracy from 0.725 to 0.8555 (gap: 0.0109). Slight overfitting emerges, but performance improves with longer training.




Figure 14: Accuracy for Momentum (100 Epochs, 30k Train/30k Validation)

2.4 Effect of Batch Size

Reducing the batch size to 32 (30,000 train/validation samples, 100 epochs) was tested.

2.4.1 Gradient Descent Optimizer

Training accuracy reaches 0.8532 and validation 0.8485 (gap: 0.0047), showing stable learning with excellent generalization despite noisier updates (Figure 15).




Figure 15: Accuracy for Gradient Descent (Batch Size 32)

2.4.2 Momentum Optimizer

Training accuracy hits 0.8740 and validation 0.8533 (gap: 0.0207), indicating higher accuracy but mild overfitting due to sensitivity to noise (Figure 16).




Figure 16: Accuracy for Momentum (Batch Size 32)

2.4.3 Comparison and Impact of Batch Size Reduction

Batch size reduction increases gradient noise. Gradient Descent maintains stability, while Momentum achieves higher accuracy with greater overfitting risk.

2.5 Performance on CPU and GPU

Time efficiency was compared over 100 epochs (batch size: 128, 50,000 training samples):

- **Gradient Descent (GPU):** 1.66 s/epoch (fastest)
- **Gradient Descent (CPU):** 1.75 s/epoch
- **Momentum (CPU):** 1.82 s/epoch
- **Momentum (GPU):** 1.83 s/epoch

GPU boosts Gradient Descent by 5.1%, but Momentum sees no benefit, making Gradient Descent with GPU the most efficient.

Optimizer	Hardware	Avg Time per Epoch (s)
Gradient Descent	GPU	1.66
Gradient Descent	CPU	1.75
Momentum	CPU	1.82
Momentum	GPU	1.83

Table 7: Time per Epoch Across Configurations

2.6 Overfitting Analysis

For the initial setup (50,000 samples, 50 epochs, batch size: 128):

- **Gradient Descent:** Training 0.8500, validation 0.8428 (gap: 0.0072)
- **Momentum:** Training 0.8618, validation 0.8521 (gap: 0.0097)

Both show minimal overfitting, with Gradient Descent slightly better at generalization.

2.7 Performance Comparison with Random Forest and SVM

Logistic regression was compared to Random Forest (100 estimators) and SVM (RBF kernel):

- **Gradient Descent:** 0.8428
- **Momentum:** 0.8521
- **Random Forest:** 0.8757
- **SVM:** 0.8790

SVM and Random Forest outperform logistic regression, with SVM leading slightly.

2.8 T-SNE Visualization of Class Weights

To further elucidate the behavior of the logistic regression model trained on the Fashion MNIST dataset, t-SNE (t-distributed Stochastic Neighbor Embedding) is employed to project the high-dimensional class weights into a two-dimensional space. This dimensionality reduction technique preserves local structures, offering a visual representation of how the Momentum and Gradient Descent optimizers distinguish the 10 clothing classes after 50 epochs of training with a batch size of 128, 50000 training set size with 10000 for validation, and a learning rate of 0.01. Separate visualizations for each optimizer reveal their effectiveness in capturing discriminative features and highlight differences in their optimization dynamics.

2.8.1 Momentum T-SNE Visualization

Figure 17 exhibits clear patterns in the learned weights, reflecting Momentum's ability to capture diverse features across categories. The spread highlights effective feature distinction, with classes like T-shirt/top and Pullover clustered closely, while Ankle boot and Sandal are positioned far apart.




Figure 17: T-SNE Visualization of Class Weights Using Momentum

2.8.2 Gradient Descent T-SNE Visualization

Figure 18 presents the t-SNE visualization for the Gradient Descent optimizer, using the same axis ranges and color coding as figure 17 for consistency. The distribution closely resembles Momentum's, with classes like Dress, Pullover, and T-shirt/top clustered together. Though the axes are wider, the weight patterns remain largely similar.

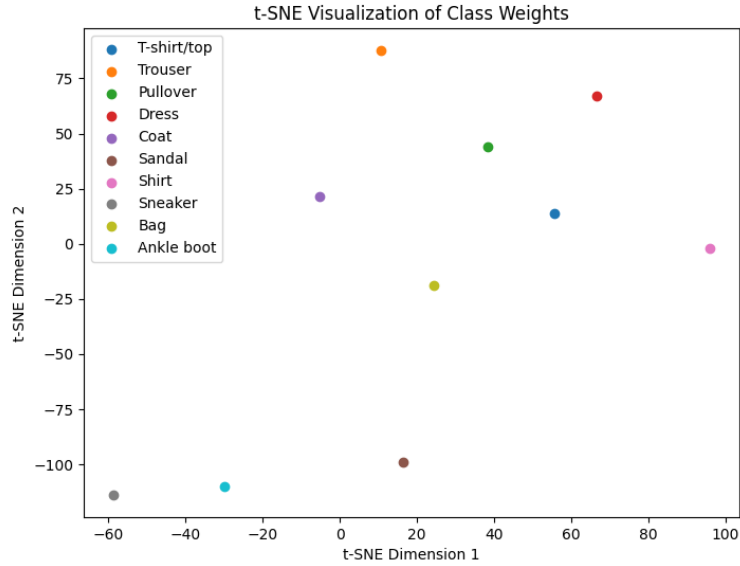


Figure 18: t-SNE Visualization of Class Weights Using Gradient Descent

References

- [1] Goodfellow, I., Bengio, Y., & Courville, A. (2016). *Deep Learning*. MIT Press.
- [2] Huber, P. J. (1964). Robust Estimation of a Location Parameter. *Annals of Mathematical Statistics*, 35(1), 73–101.